

Contextual String Embeddings for Sequence Labeling

Alan Akbik
zalando Research
Mhlenstrae 25
10243 Berlin

Alan Akbik@zalando.de

Duncan Blythe
zalando Research
Mhlenstrae 25
10243 Berlin

Duncan Blythe@zalando.de

Roland Vollgraf
zalando Research
Mhlenstrae 25
10243 Berlin

Riland Vollgraf@zalando.de

1. Introduction

As a basic task and foundation work of NLP, sequence labeling has a significant influence on the follow-up works, such as named entity recognition (NER), part-of-speech (PoS) tagging, etc. For the traditional machine learning methods, we initially use HMM (Hidden Markov Model) to solve the problem. Regarded the tag sequence as hidden state sequence and sentence as an obvious sequence, HMM calculates the probability distribution of each moment by defining the 5-tuple flow (hidden state set S , the observed value set O , state transition matrix A , initial state probability distribution π , observed value probability distribution of specified state B). Generally speaking, HMM involves two basic tasks, assessment and decoding. The assessment aims to calculate the probability of the observed sequence, and it always adopts a backward or forward algorithm. While decoding is required to ensure the most likely hidden state sequence based on the obvious sequence, which is selects Viterbi algorithm. However, because HMM determines each hidden states individually, rather than considering the dependency of interlink states, the result that labeling two or more prepositions simultaneously will happen in the part-of-speech tagging. To overcome this shortcoming, the simple way is to add the connection matrix at adjacent hidden states, therefore, by improving maximum the entropy model, the CRF (Conditional random field) comes into being. Relying on hand-feature descriptions, CRF applied widely at pre-deep learning era. With deep neural network rapidly developed, depending on the powerful feature representation ability, BiLSTM+CRF took replace to CRF and renewed the record of sequence labeling.

About embeddings, the original form is one-hot discrete word vectors. Although this method is simplified, it exists a semantic gap and the curse of dimensionality. After that, the word bag model (counting the times of words in the files), TF-IDF (reflecting the proportion of the words in files). Based on the distributional hypothesis (the words have the similar means with the same context), Bengio et al. (2001)

invented NNLM (Neural Network Language Model) and presented the word vector, which converts the one-hot code to the dense vectors by looking-up operation and the look-up table is created by the training process. According to the input distinguish, word vector networks include NNLM, C&W, CBOW, Skip-gram, etc. Comparing to the NNLM, C&W as well as Skip-gram delete the hidden layer and ignore the order of the words, making the architecture slender. Besides, for the reason that the softmax calculation needs considerable computing resource and time, the researcher proposed hierarchical softmax, constructing the Huffman tree to update the vocabulary partially. Furthermore, about the polysemy like "bank", we can embedding multi-word vectors for different contexts, then using k-means clustering choose appropriate word vectors. Some experiments show concatenating different word vectors training by different network architectures, such as RNN, CNN, LSTM, can also improve the F1 value.

The biggest innovation of this paper about the sequence labeling task is to improve the tradition embedding, the author propose to leverage the internal states of a trained character language model to produce a novel type of word embedding which referring to as contextual string embeddings. These embeddings have the distinct properties that they (a) are trained without any explicit notion of words and thus fundamentally model words as sequences of characters, and (b) are contextualized by their surrounding text, meaning that the same word will have different embeddings depending on its contextual use.

2. Major contributions and the proposed solutions

Figure1. is the detailed structure of CLM, use characters as atomic units of language modeling, allowing text to be treated as a sequence of characters passed to an LSTM which at each point in the sequence is trained to predict the next character¹. This means that the model possesses a hidden state for each character in the sequence.

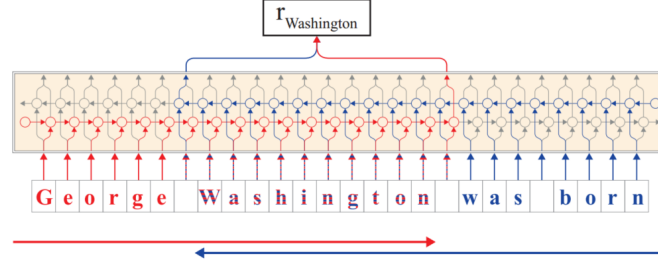


Figure 1. Extraction of a contextual string embedding for a word (Washington) in a sentential context. From the forward language model (shown in red), we extract the output hidden state after the last character in the word. This hidden state thus contains information propagated from the beginning of the sentence up to this point. From the backward language model (shown in blue), we extract the output hidden state before the first character in the word. It thus contains information propagated from the end of the sentence to this point. Both output hidden states are concatenated to form the final embedding.

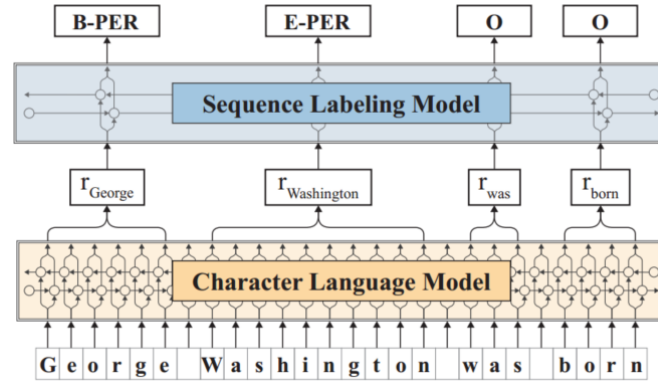


Figure 2. High level overview of proposed approach. A sentence is input as a character sequence into a pre-trained bidirectional character language model. From this LM, we retrieve for each word a contextual embedding that we pass into a vanilla BiLSTM-CRF sequence labeler, achieving robust state-of-the-art results on downstream tasks (NER in Figure).

Figure2. is the sequence labeling model detail. The network can split two parts, Character Language Model (CLM) and Sequence Labeling Model (SLM). CLM leverages BiLSTM hidden state encoding word vectors, on account of BiLSTM merges the context information naturally, thus it can produce different embeddings for polysemous words depending on their usage and model words and context fundamentally as sequences of characters, to both better handle rare and misspelled words as well as model subword structures such as prefixes and endings. The CLM can be utilized to downstream sequence labeling tasks effectively. As for name entity recognition (NER), the author utilizes BiLSTM-CRF model.

Whatsmore, the author evaluate the proposed contextual string embedding in the five configurations: Proposed, Proposed+word, Proposed+char, Proposed+word+char, Proposed+all. The experiment shows that the added use of classic word embeddings in setup Proposed+word often produces the best results. Especially for NER, the use of classic word embeddings increases average F1 score by 1.1 pp to 93.07, compared with 91.97 for Proposed.

3. Own understanding of the contributions and the unsolved problems of the paper

Comprehensively, for the sequence labeling tasks, there are two mainstream methods currently, the first one is based on the abundant high-quality corpus to pre-training word vectors, like Bert, and the second one is based on the feature representation and selection ability of networks to extract the hidden state and embedding or concatenating various word vectors so as to improve results.

4. Detailed Information

@inproceedingsakbik2018contextual,
 title=Contextual string embeddings for sequence labeling,
 author=Akbik, Alan and Blythe, Duncan and Vollgraf, Roland,
 booktitle=Proceedings of the 27th International Conference on Computational Linguistics,
 pages=1638–1649,
 year=2018